

# Homework 3: Autoencoders

Hai Dinh (19960331-4494)

## 1. Presentation of results

Figure 1 and Figure 2 below show the training progress for autoencoder 1 and autoencoder 2. As can be seen, the training performance sharply increases within only the first 50 epochs or so, and then quickly levels off as the training goes on. Validation patience can actually be used here to reduce the training time by stopping early. However, since the training goes very quickly on my GPU, so I decided to run the network all the way until epoch 800.

Autoencoder 1 reaches approximately the final loss of 0.067, while autoencoder 2 got the lowest loss of 0.049. Clearly, autoencoder 2 performs better, which is no surprise because autoencoder 2 has more hidden units in the bottleneck layer, thus is able to include more information.




Figure 1: Training performance of autoencoder 1





Figure 2: Training performance of autoencoder 2

Montage (inputs/model1/model2)



*Figure 3: Montage where first column is the inputs, second column is the outputs of autoencoder 1, and third column is the outputs of autoencoder 2.*

As can be seen here in the montage in Figure 3:

- Autoencoder 1 seems to be able to predict digit 1 very well. Digits 0 and 9 can be recognized a bit, if we don't judge harshly. The rest of the digits are really bad.
- Autoencoder 2 seems to perform well on digits 0, 1, 3, and 9. Other digits are not so bad either, as the forms of the digits can be recognized.

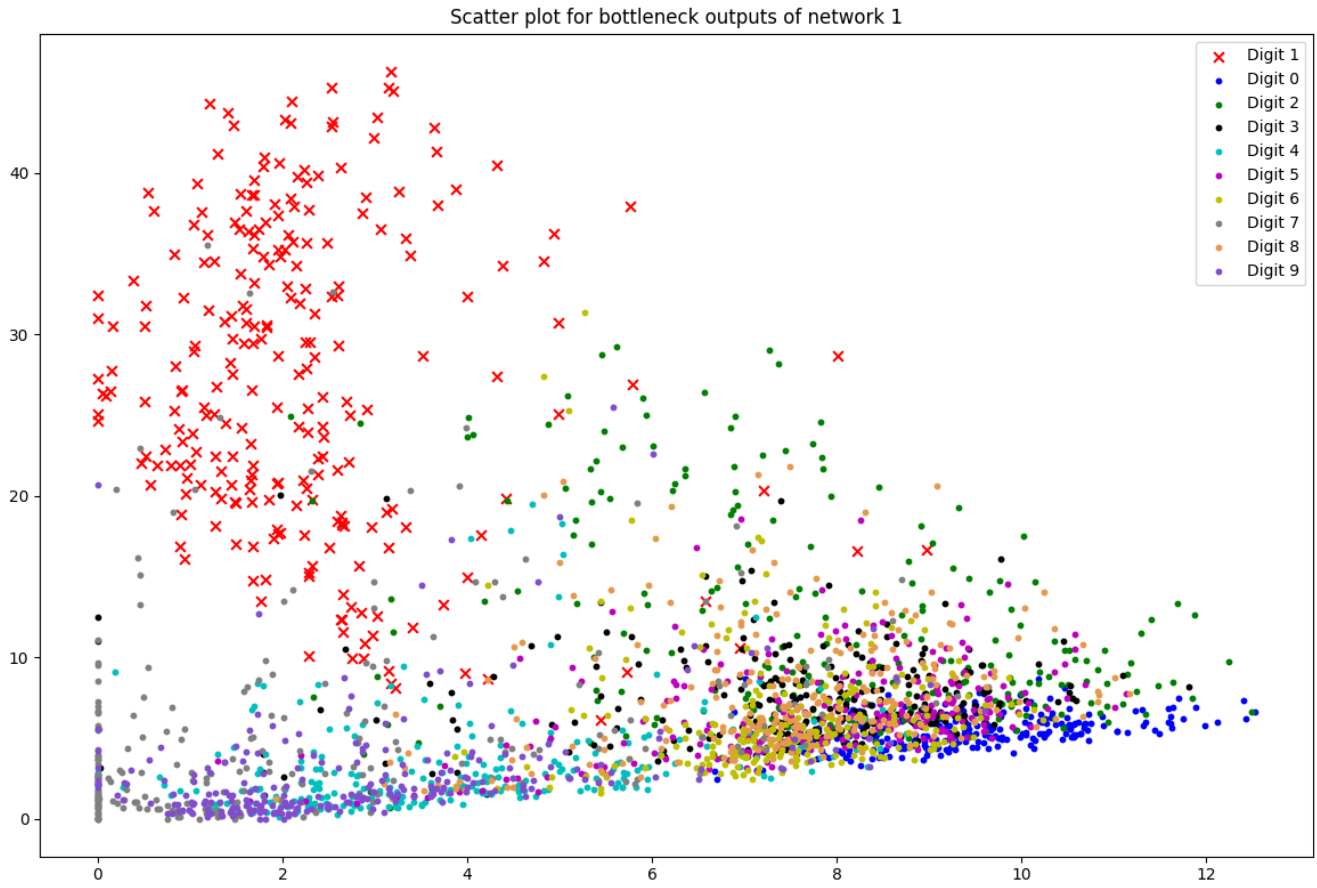


Figure 4: Scatter plot showing the outputs of the bottleneck layers for autoencoder 1. Here, the X markers represent well-recognized digits and O markers represent the non-well-recognized digits.

As shown in the scatter plot in Figure 3:

- Digit 1 is the most well-recognized digit, so I plotted that with X marker.
- Even though digit 0 and digit 9 are plotted with O marker, they are a little bit better than the other digits. However, they can't be called as "well-recognized", since digit 9 can still be mistaken as digit 7, and digit 0 can still be mistaken as other digits (at least it overlaps quite a bit with the plots of other digits).

From the scatter plot (Figure 4), it looks like the upper left corner is the region that most likely belongs to digit 1. The lower left corner is the region for digit 9 (though it can also be mistaken as digit 7). And the lower right corner is the region for digit 0 (though it can also be mistaken for other digits in the same region as well). To test these “rules”, I created 3 points in 2D at these corners. Then I simply passed these 3 inputs into the decoder, and to see if the decoder can generate the expected digits.

- For digit 0, I chose the point (12, 4) in the right bottom corner.
- For digit 1, I chose the point (2, 40) in the upper left corner.
- For digit 9, I chose the point (1.5, 0) in the left bottom corner.




Figure 5: Outputs from the decoder of autoencoder 1, given the points (12, 4) and (2, 40) and (1.5, 0) as inputs to the decoder.

To find the coding rule for autoencoder 2, let’s compute the mean output vectors of the encoder for each of the digits. Here is the result:

```
[INFO] Mean bottleneck vectors for each digit:
Digit 0 ==> [10.978588 12.026303 6.9547396 8.330039 ]
Digit 1 ==> [37.925915 15.457135 71.48261 12.837233]
Digit 2 ==> [6.7864785 7.5168977 6.0474663 3.2689466]
Digit 3 ==> [14.517906 12.474363 10.927933 7.5871468]
Digit 4 ==> [ 7.212148 6.836543 12.227403 13.035679]
Digit 5 ==> [11.919395 12.632617 11.724218 10.654239]
Digit 6 ==> [5.9780188 6.209823 5.099382 4.002297 ]
Digit 7 ==> [12.8656845 12.9306755 30.59925 27.417377 ]
Digit 8 ==> [11.773857 12.73176 13.145078 10.882374]
Digit 9 ==> [10.887202 9.4659815 18.285233 18.306738 ]
```

**Model2: Plot of decoder outputs for mean bottleneck vectors for each digit**




Figure 6: Outputs from the decoder of autencoder 2, given the mean bottleneck vectors as inputs for each of the digits. Here digits 0, 1, 9 are classified pretty well. Digits 3 and 7 are also quite good (though it’s a bit blurry, and 7 looks a bit like a 9).

By inspection of the mean bottleneck vectors above, I managed to work out a couple rules of thumb, at least for the well-classified digits such as 0, 1, and 9.

- For digit 0: The first 2 dimensions are big, the last 2 only half as big. The 2nd dimension is slightly bigger than the 1st dimension, and the 4th dimension is slightly bigger than the 3rd dimension. To fit this criteria, I have chosen this point: (50, 60, 30, 40).
- For digit 1: The 1st and 3rd dimensions are big, the 2nd and 4th dimensions are small. The 1st dimension is approximately half as big as the 3rd dimension. To fit this criteria, I have chosen this point: (50, 10, 90, 10).
- For digit 9: The first 2 dimensions are approximately half as big as the last 2 dimensions. To fit this criteria, I have chosen this point: (50, 50, 90, 90).

Passing these points to the decoder, we get the outputs as in Figure 7. All 3 digits are classified correctly, which means the rules of thumb that I found worked quite well.



*Figure 7: Outputs from the decoder of autencoder 2, given the points (50, 60, 30, 40) and (50, 10, 90, 10) and (50, 50, 90, 90) as inputs.*

## 2. Code

```
import os
import argparse
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras.layers import Dense, InputLayer

def load_mnist(val_seed=None):
    """
    Load, pre-process, and split the MNIST dataset into train/val/test sets.

    Arguments:
        val_seed (int): Seed to generate the validation set.

    Returns:
        (x_train, y_train): ndarray of shape (50000, 784) and (50000,) representing train set.
        (x_val, y_val)      : ndarray of shape (10000, 784) and (10000,) representing val set.
        (x_test, y_test)   : ndarray of shape (10000, 784) and (10000,) representing test set.
    """
    saved_random_generator_state = np.random.get_state()
    np.random.seed(val_seed)

    (x, y), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
    x, x_test = x.reshape(60000, 784), x_test.reshape(10000, 784)
    x, x_test = x / 255.0, x_test / 255.0

    indices = np.random.permutation(x.shape[0])
    train_indices, val_indices = indices[:50000], indices[50000:]
    x_train, y_train = x[train_indices], y[train_indices]
    x_val, y_val = x[val_indices], y[val_indices]

    print("[INFO] Train set shapes:", x_train.shape, y_train.shape)
    print("[INFO] Validation set shapes:", x_val.shape, y_val.shape)
    print("[INFO] Test set shapes:", x_test.shape, y_test.shape)

    np.random.set_state(saved_random_generator_state)
    return (x_train, y_train), (x_val, y_val), (x_test, y_test)

def construct_network(bottleneck_size):
    """
    Construct the autoencoder network.

    Arguments:
        bottleneck_size (int): Number of neurons in the bottleneck layer.

    Returns:
        model: The model containing the entire autoencoder network.
    """
    model = tf.keras.models.Sequential([
        InputLayer(input_shape=(784)),
        Dense(50, kernel_initializer="glorot_uniform", activation="relu"),
        Dense(bottleneck_size, kernel_initializer="glorot_uniform", activation="relu"),
        Dense(784, kernel_initializer="glorot_uniform", activation="relu"),
    ])

```

```

model.compile(
    optimizer=tf.optimizers.Adam(learning_rate=0.001),
    loss="mean_squared_error",
)
model.summary()
return model

def split_network(model):
    """
    Split the network into 2 parts (encoder and decoder)

    Arguments:
        model (Sequential): The sequential model to be split.

    Returns:
        The encoder and decoder as 2 separate networks.
    """
    first_layer = model.get_layer(index=0)
    bottleneck_layer = model.get_layer(index=1)
    last_layer = model.get_layer(index=2)

    encoder = tf.keras.models.Sequential([
        InputLayer(input_shape=(784)),
        first_layer,
        bottleneck_layer,
    ])
    decoder = tf.keras.models.Sequential([
        InputLayer(input_shape=(bottleneck_layer.units)),
        last_layer,
    ])
    return encoder, decoder

def train_network(mnist, network, outdir):
    """
    Train the autoencoder network.

    Arguments:
        network (int): Either 1 or 2, indicating which autoencoder network to train.
        outdir (str): Directory to save the output to.
    """
    print("\n[INFO] Training network {}".format(network))
    (x_train, y_train), (x_val, y_val), (x_test, y_test) = mnist
    model = construct_network(2 if network == 1 else 4)

    training_performance = model.fit(
        x_train,
        x_train,
        validation_data=(x_val, x_val),
        shuffle=True,
        batch_size=8192,
        epochs=800,
        verbose=2,
    )
    training_performance = training_performance.history
    plt.figure()
    plt.title("Autoencoder {}".format(network))
    plt.plot(training_performance["loss"])

```



```

plt.plot(training_performance["val_loss"])
plt.legend(['Train', 'Validation'])
plt.ylabel("Mean squared error")
plt.xlabel("Epoch")
plt.savefig("{}autoencoder{}.png".format(outdir, network))
model.save("{}autoencoder{}.h5".format(outdir, network))

def load_network(network, outdir):
    """
    Load a trained autoencoder network.

    Arguments:
        network (int): Either 1 or 2, indicating which autoencoder network to train.
        outdir (str): Directory to load the model from.

    Returns:
        The trained model.
    """
    print("\n[INFO] Loading model from {}autoencoder{}.h5".format(outdir, network))
    return tf.keras.models.load_model("{}autoencoder{}.h5".format(outdir, network))

def create_montage(mnist, model1, model2, outdir, test_seed=None):
    """
    Create montage to compare results between model1 and model2 on random samples from the test set.

    Args:
        mnist: The MNIST dataset as loaded from load_mnist().
        model1: The loaded model for autoencoder 1.
        model2: The loaded model for autoencoder 2.
        outdir: Directory to save the montage to.
        test_seed: Seed to randomly sample 10 digits from the test set.
    """
    (x_train, y_train), (x_val, y_val), (x_test, y_test) = mnist
    saved_random_generator_state = np.random.get_state()
    np.random.seed(test_seed)
    indices = [np.random.choice(np.where(y_test == i)[0]) for i in range(10)]
    np.random.set_state(saved_random_generator_state)

    inputs = x_test[indices].reshape(280, 28)
    output1 = model1.predict(x_test[indices]).reshape(280, 28)
    output2 = model2.predict(x_test[indices]).reshape(280, 28)
    montage = np.concatenate([inputs, output1, output2], axis=1)

    plt.figure()
    plt.title("Montage (inputs/model1/model2)")
    plt.imshow(montage, cmap="gray_r")
    plt.axis("off")
    plt.savefig("{}montage.png".format(outdir))

def create_scatter(mnist, encoder, decoder, outdir):
    """
    Create scatter plot for outputs of the bottleneck layer of the first autoencoder.

    Arguments:
        mnist: The MNIST dataset as loaded from load_mnist().
        encoder: Encoder of the first autoencoder.

```

```

    decoder: Decoder of the first autoencoder.
    outdir: Directory to save the scatter plot to.
"""
(x_train, y_train), (x_val, y_val), (x_test, y_test) = mnist

def scatter(digit, color, n, well_recognized):
    indices = np.where(y_test == digit)[0][:250]
    outputs = encoder.predict(x_test[indices])
    marker = "x" if well_recognized else "."
    label = "Digit " + str(digit)
    plt.scatter(outputs[:,0], outputs[:,1], c=color, marker=marker, s=40, label=label)

good_digits = [1]
good_colors = ["r"]
bad_digits = [0, 2, 3, 4, 5, 6, 7, 8, 9]
bad_colors = ["b", "g", "k", "c", "m", "y", [(0.5, 0.5, 0.5)], [(0.9, 0.6, 0.3)], [(0.5, 0.3, 0.8)]]

plt.figure(figsize=(12,8))
[scatter(d, c, 100, well_recognized=True) for d, c in zip(good_digits, good_colors)]
[scatter(d, c, 100, well_recognized=False) for d, c in zip(bad_digits, bad_colors)]
plt.legend()
plt.tight_layout(pad=2)
plt.title("Scatter plot for bottleneck outputs of network 1")
plt.savefig("{}scatter.png".format(outdir))

def do_experiment_on_model1_rules(decoder, outdir):
    """
    Find the rules for the first autoencoder. Here we are gonna try to create the digit 1 (which is
    well-recognized), and digits 0/9 (which are somewhat well-recognized), by only using the decoder
    of the first autoencoder. The idea is to use the coordinates that are "well-separated" from the
    scatter plots, and see if the decoder will give the expected digits or not.
    """
    fig = plt.figure(figsize=(6,3))
    plt.title("Rules for autoencoder 1")
    plt.axis("off")

    def generate(counter, digit, bottleneck_point):
        fig.add_subplot(1, 3, counter)
        plt.axis("off")
        plt.title("Expected digit: " + str(digit))
        output = decoder.predict(bottleneck_point).reshape(28,28)
        plt.tight_layout(pad=0)
        plt.imshow(output, cmap="gray_r")

    generate(1, digit=0, bottleneck_point=np.array([[12, 4]]))
    generate(2, digit=1, bottleneck_point=np.array([[2, 40]]))
    generate(3, digit=9, bottleneck_point=np.array([[1.5, 0]]))
    plt.savefig(outdir + "/rules_model1.png")

def do_experiment_on_model2_rules(mnist, encoder, decoder, outdir):
    """
    Find the rules for the second autoencoder. Here we are gonna try to find the mean bottleneck
    output vectors for each digit, and then feed those mean vectors to the decoder to see if we get
    the expected digits or not.
    """
    (x_test, y_test), (x_val, y_val), (x_test1, y_test1) = mnist
    bottleneck_vectors = encoder.predict(x_test)

```

```

mean_vectors = [np.mean(bottleneck_vectors[y_test == digit], axis=0) for digit in range(10)]

print("\n[INFO] Mean bottleneck vectors for each digit:")
for digit, mean in enumerate(mean_vectors):
    print("Digit {} ==> {}".format(digit, mean))

decoder_inputs = np.array(mean_vectors)
decoder_outputs = decoder.predict(decoder_inputs)
decoder_outputs = np.concatenate(decoder_outputs.reshape(10,28,28), axis=1)

plt.figure(figsize=(8, 1.5))
plt.title("Model2: Plot of decoder outputs for mean bottleneck vectors for each digit")
plt.imshow(decoder_outputs, cmap="gray_r")
plt.axis("off")
plt.savefig("{} /mean_vector_outputs.png".format(outdir))

fig = plt.figure(figsize=(6,3))
plt.title("Rules for autoencoder 2")
plt.axis("off")

def generate(counter, digit, bottleneck_point):
    fig.add_subplot(1, 3, counter)
    plt.axis("off")
    plt.title("Expected digit: " + str(digit))
    output = decoder.predict(bottleneck_point).reshape(28,28)
    plt.tight_layout(pad=0)
    plt.imshow(output, cmap="gray_r")

generate(1, digit=0, bottleneck_point=np.array([[50, 60, 30, 40]]))
generate(2, digit=1, bottleneck_point=np.array([[50, 10, 90, 10]]))
generate(3, digit=9, bottleneck_point=np.array([[50, 50, 90, 90]]))
plt.savefig(outdir + "/rules_model2.png")

def main(args):
    """
    Main program to train and evaluate autoencoders.
    """
    mnist = load_mnist(val_seed=123)

    if not args.no_training:
        train_network(mnist, 1, args.outdir)
        train_network(mnist, 2, args.outdir)

    model1 = load_network(1, args.outdir)
    model2 = load_network(2, args.outdir)
    encoder1, decoder1 = split_network(model1)
    encoder2, decoder2 = split_network(model2)

    create_montage(mnist, model1, model2, args.outdir)
    create_scatter(mnist, encoder1, decoder1, args.outdir)
    do_experiment_on_model1_rules(decoder1, args.outdir)
    do_experiment_on_model2_rules(mnist, encoder2, decoder2, args.outdir)

if __name__ == "__main__":
    parser = argparse.ArgumentParser(description="Autoencoder networks for MNIST")
    parser.add_argument("--no-training", "-nt", action="store_true", help="Perform no training")
    parser.add_argument("--outdir", "-o", type=str, default=".", help="Out directory")

```

```

args = parser.parse_args()

os.environ["TF_CPP_MIN_LOG_LEVEL"] = "2" # remove info/warning logs from tensorflow
gpus = tf.config.experimental.list_physical_devices("GPU")

if len(gpus) > 0:
    try:
        tf.config.experimental.set_memory_growth(gpus[0], True)
    except RuntimeError as e:
        print(e)

print("[INFO] Using Tensorflow version:", tf.__version__)
print("[INFO] Number of GPUs available:", len(gpus))
main(args)

```